

VOLUME III

AI POWER USER PROMPTS

Advanced prompt categories for serious AI practitioners

System Prompt Design · Memory & Context · Data & Spreadsheets · Prompt Chaining Negotiation & Persuasion ·
Legal & Compliance · Education Design · Productivity Social Media & Community · AI Image Editing & Inpainting

10 ADVANCED CATEGORIES · 40 POWER PROMPTS

Completes the trilogy with Vol. I: Best Proven AI Prompts & Vol. II: AI Tool Master Prompts

SYSTEM PROMPT & INSTRUCTION DESIGN

POWER PROMPT	WHAT IT EFFECTIVELY DOES
› You are [PERSONA NAME], a [role description]. Your core purpose is [primary goal]. You always [non-negotiable behaviour 1] and [non-negotiable behaviour 2]. You never [hard constraint 1] or [hard constraint 2]. When uncertain, you [fallback behaviour]. Your tone is [tone adjectives].	The master system prompt template. Each clause handles a specific failure mode: persona prevents drift, non-negotiables catch edge cases, the fallback prevents hallucination under uncertainty, and tone anchors every response.
› Design a system prompt for an AI assistant that handles [use case]. It should: excel at [primary task], gracefully decline [out-of-scope requests], always ask for clarification when [trigger conditions], and maintain [specific persona traits]. Test it against these edge cases: [list 3 edge cases].	System prompt design with built-in QA. The edge case testing step catches gaps in the instruction set before deployment — most custom GPT failures stem from untested edge cases.
› Audit this system prompt: [SYSTEM PROMPT]. Identify: instructions that conflict with each other, missing constraints that could cause harmful outputs, ambiguous language the model might misinterpret, and gaps where the model has no guidance. Rewrite an improved version.	System prompt red-teaming. Conflicting instructions are the most common source of inconsistent AI behaviour — this audit surfaces them before users encounter them.
› Write a system prompt that makes the AI respond only in [specific format/structure] for every response, regardless of how the user asks. Include instructions for handling requests that cannot fit the format, and a graceful override command the user can invoke.	Format-locked system design with an escape hatch. The override command is critical — without it, rigid format instructions break down on legitimate exceptional requests.

MEMORY & CONTEXT MANAGEMENT

POWER PROMPT	WHAT IT EFFECTIVELY DOES
› We are starting a long project together. At the end of each conversation, generate a CONTEXT HANDOFF in this format: [KEY DECISIONS MADE], [OPEN QUESTIONS], [NEXT ACTIONS], [IMPORTANT CONSTRAINTS WE ESTABLISHED], [WHAT TO REMIND ME OF NEXT SESSION]. I will paste this at the start of our next conversation.	Stateless memory architecture using structured handoff. Solves the context window problem without external memory tools. The 'remind me' field catches follow-ups that would otherwise fall through the cracks.
› Before we continue, here is the context from our last session: [HANDOFF BLOCK]. Confirm you have understood it by summarising the 3 most important constraints and decisions in your own words before we proceed.	Context verification prompt. The confirmation step catches misunderstood context before it corrupts subsequent work — critical when handoff blocks are long or complex.
› This conversation is getting long. Compress everything we have discussed into a dense summary that preserves: all decisions made, all constraints established, all code/content produced (summarised), and all unresolved questions. Make it short enough to paste at the top of a new conversation.	Mid-session context compression. Prevents degraded performance as conversations approach token limits by proactively distilling before quality drops.

› Maintain a running [DECISION LOG] throughout our conversation. Every time we make a decision, choose an approach, or establish a constraint, add it to the log. Show me the current log on request or whenever it feels relevant.

In-session memory scaffolding. The decision log prevents the common failure of contradicting earlier choices in long sessions — especially valuable in design and strategy conversations.

DATA & SPREADSHEET AI

POWER PROMPT	WHAT IT EFFECTIVELY DOES
› Analyse this dataset: [DATA]. Give me: a data quality report (missing values, outliers, inconsistencies), the 5 most statistically significant patterns, a hypothesis for why each pattern exists, and the 3 most actionable insights a non-technical stakeholder should act on.	End-to-end data analysis in one prompt. The quality report step catches data issues that would invalidate insights. The non-technical reframe makes outputs usable by decision-makers, not just analysts.
› Clean this messy data: [DATA]. Tasks: standardise date formats to [format], fix inconsistent category labels, remove duplicates, flag (don't delete) rows with suspicious values, and explain every change you made in a changelog.	Auditable data cleaning. The 'flag, don't delete' instruction preserves data integrity. The changelog is critical for regulatory compliance and reproducibility in professional data work.
› Write an Excel/Google Sheets formula to: [task description]. The formula should handle: empty cells without breaking, [edge case 1], and [edge case 2]. Explain each component of the formula in plain English.	Formula design with failure-mode handling baked in. Empty cell handling is the most common formula failure in production spreadsheets — this prompt makes it mandatory.
› I have this data: [DATA]. Suggest 5 different ways to visualise it, ranked by how clearly each communicates [the key insight I want to show]. For the top choice, describe the exact chart type, axes, colour coding, and what annotation to add.	Visualisation strategy before execution. The insight-first framing prevents the common failure of choosing a chart type before knowing what story the data should tell.

PROMPT CHAINING & WORKFLOWS

POWER PROMPT	WHAT IT EFFECTIVELY DOES
› I am building a multi-step AI workflow. Step 1 output will be the input for Step 2, and so on. Design the complete chain for: [GOAL]. For each step specify: the exact prompt, the expected output format, validation checks before passing to next step, and what to do if a step fails.	Full workflow architecture with error handling. The validation checks between steps are what separate robust pipelines from fragile ones — most prompt chains fail silently at handoff points.
› Take the output from this previous step: [STEP 1 OUTPUT]. Your job in this step is to: [STEP 2 TASK]. Format your output as [FORMAT] so it can be directly used by the next step. Flag anything in the input that seems like an error from the previous step.	The inter-step handoff prompt. The error-flagging instruction creates self-healing pipelines where bad upstream outputs are caught rather than silently propagated.

› You are the quality control step in a multi-step AI pipeline. Review this output: [OUTPUT]. Check it against the original brief: [BRIEF]. Score it 1-10 on accuracy, completeness, and format compliance. If below 8 on any dimension, return it with specific revision instructions. If 8+, approve it and summarise what was delivered.	Automated QA gate in AI workflows. The scored approval/rejection loop with specific revision instructions creates a reviewable audit trail and catches low-quality outputs before they propagate.
› Design a prompt chain that takes [RAW INPUT] and produces [FINAL OUTPUT] in no more than 4 steps. For each step: name it, write the prompt, specify input/output format, and estimate the token cost. Optimise for reliability over speed.	Constrained chain design. The 4-step limit forces simplicity — the most reliable pipelines are the shortest ones. The token cost estimate makes scaling costs visible before building.

NEGOTIATION & PERSUASION

POWER PROMPT	WHAT IT EFFECTIVELY DOES
› I am negotiating [situation — salary, contract, price, terms]. My position: [MY POSITION]. Their likely position: [THEIR POSITION]. My BATNA: [BEST ALTERNATIVE]. Give me: an opening move, 3 concession sequences I can offer, responses to their 3 most likely objections, and a walk-away signal to watch for.	Complete negotiation playbook. The BATNA input is the key variable — it calibrates how aggressive the strategy should be. The walk-away signal prevents over-negotiating a deal that was already won.
› Write a persuasive case for [POSITION] targeting [SPECIFIC DECISION MAKER]. Research their known priorities and concerns: [WHAT YOU KNOW ABOUT THEM]. Lead with what they care about most, not what I care about most. Anticipate and pre-empt their top 2 objections.	Audience-first persuasion. The 'lead with what they care about' constraint flips the common mistake of leading with your own interests. Pre-empting objections is significantly more effective than responding to them.
› I need to have a difficult conversation with [PERSON] about [ISSUE]. Draft what I should say. Tone: direct but not aggressive. Acknowledge their perspective before stating mine. Have a clear ask at the end. Also give me 3 ways they might respond and how I should handle each.	Difficult conversation scripting with response trees. The acknowledgement-before-position structure is the single most effective de-escalation technique. Response trees prevent being caught off-guard.
› Analyse this email/message I received: [MESSAGE]. What is the sender really asking for beneath the surface? What emotions or concerns are driving it? What is the most strategic response that addresses both the stated and unstated need? Draft that response.	Subtext analysis before responding. Reading unstated needs is the difference between a response that closes a conversation and one that escalates it. Especially powerful for customer complaints and stakeholder management.

LEGAL & COMPLIANCE

POWER PROMPT	WHAT IT EFFECTIVELY DOES
› Review this contract/document: [DOCUMENT]. Identify: clauses that are unusually unfavourable to me, missing standard protections, ambiguous language that could be interpreted against my interests, termination conditions I should understand, and the top 3 things I should negotiate before signing. Note: this is not legal advice.	Pre-lawyer contract review that surfaces the issues worth paying a lawyer to fix. The 'top 3 to negotiate' distillation prioritises review time — most people miss critical clauses by reading linearly.

› Translate this legal/regulatory document into plain English: [DOCUMENT]. For each section: state what it requires in one sentence, who it applies to, what happens if you don't comply, and flag any section that is genuinely ambiguous or requires expert interpretation.	Compliance translation with ambiguity flags. The 'requires expert interpretation' flag is essential — it prevents dangerous over-confidence in AI-interpreted legal text.
› Draft a [contract type / policy / terms of service] for [use case]. Include standard protections for [my role]. Flag any clause where the law varies significantly by jurisdiction. Write in plain language — no legalese unless legally required.	Plain-language legal drafting with jurisdiction flags. The jurisdiction flag is the most important guard rail — legal requirements vary dramatically by location and this prompt makes that visible.
› I need to comply with [REGULATION — GDPR, HIPAA, SOC2, etc.]. Audit this [policy / process / system description] against the regulation's requirements. List: what we are doing right, what we are missing, what is ambiguous, and the highest-risk gaps to address first.	Regulatory gap analysis with risk prioritisation. The 'highest-risk gaps first' instruction is what makes this actionable — compliance remediation always has resource constraints and needs triage.

EDUCATION & LEARNING DESIGN

POWER PROMPT	WHAT IT EFFECTIVELY DOES
› Design a complete curriculum for teaching [TOPIC] to [AUDIENCE] in [TIMEFRAME]. Include: learning objectives (what they can DO, not just know), module sequence with rationale, one hands-on project per module, assessment criteria, and common misconceptions to address proactively.	Outcomes-based curriculum design. The 'can DO, not just know' instruction shifts from passive knowledge transfer to active skill development — the difference between a course people complete and one that changes behaviour.
› Explain [CONCEPT] to me as if I am [beginner / intermediate / expert]. Use one analogy, one real-world example, and one common mistake people make when first learning this. Then give me 3 questions I should be able to answer if I truly understand it.	Level-calibrated explanation with built-in comprehension check. The misconception instruction proactively closes the gaps that cause learners to think they understand when they don't.
› Create a spaced repetition study plan for [TOPIC]. Generate: 20 flashcard questions covering core concepts (with answers), 5 application questions requiring synthesis, and a review schedule for the next 30 days. Identify the 3 concepts most likely to be confused with each other.	Evidence-based study system in one prompt. The confusion-mapping instruction at the end targets the most common comprehension failures — these are the concepts worth drilling hardest.
› I want to learn [SKILL] as fast as possible. I have [TIME AVAILABLE] per day and [WEEKS] to learn it. Design a minimum effective dose learning plan: what to learn, in what order, and what to skip entirely. Include one deliberate practice exercise per week.	Pareto learning design. The 'what to skip' instruction is the most valuable part — it forces the model to identify the 20% of content that delivers 80% of competency, which most curricula never do.

PERSONAL PRODUCTIVITY & LIFE

POWER PROMPT	WHAT IT EFFECTIVELY DOES
--------------	--------------------------

› Act as my weekly review facilitator. Ask me these questions one at a time and help me reflect deeply: What did I accomplish this week? What did I not finish and why? What drained my energy? What energised me? What is my top priority next week? What am I avoiding? Then summarise my answers into 3 action commitments.	Structured weekly review with coaching. The one-at-a-time delivery prevents overwhelm and creates genuine reflection. The 'what are you avoiding' question is the one most people skip and most need.
› Help me make this decision: [DECISION]. I will give you my options: [OPTIONS]. Ask me 5 questions that will surface my true values and priorities relevant to this decision. Then, based on my answers, give me a recommendation — don't hedge it.	Values-elicitation before recommendation. The questioning step externalises the real criteria before analysis. The 'don't hedge' instruction prevents the useless 'it depends on your priorities' non-answer.
› I want to build the habit of [HABIT]. My current situation: [CONTEXT]. Design a minimum viable habit that I could not reasonably fail at, a trigger to attach it to, a reward, and a way to track it. Then design 3 monthly progressions to scale it over 90 days.	Habit design rooted in behavioural science. The 'could not reasonably fail' constraint forces starting absurdly small — the evidence-based key to habit formation. The 90-day progression prevents the plateau that kills habits at week 3.
› I have these tasks: [TASK LIST]. Help me prioritise using the Eisenhower Matrix. Then reorder my day with time blocks, protecting [my most important work] in my peak energy hours ([TIME RANGE]). Flag anything I should delegate or delete entirely.	Prioritisation with energy-aware scheduling. The peak energy hours input is what most productivity systems miss — doing the right task at the wrong time is nearly as bad as doing the wrong task.

SOCIAL MEDIA & COMMUNITY

POWER PROMPT	WHAT IT EFFECTIVELY DOES
› Write a LinkedIn post about [TOPIC/STORY]. Format: hook (1 provocative line), context (2-3 lines), insight (the non-obvious point), evidence (one specific example or data point), takeaway (one actionable line), and soft CTA. No hashtag spam. Max 1,300 characters.	LinkedIn post architecture optimised for the algorithm and for humans. The hook as a single provocative line drives the 'see more' click. The no-hashtag-spam instruction reflects how the LinkedIn algorithm actually ranks content.
› Write a Twitter/X thread on [TOPIC]. Thread structure: tweet 1 is a bold standalone claim, tweets 2-8 each add one specific piece of evidence or nuance, tweet 9 is the synthesis, tweet 10 is the provocative question or CTA. Each tweet max 260 characters. No filler tweets.	Thread architecture that earns retweets. The bold standalone claim means tweet 1 works as a quote tweet or standalone share. The 'no filler' instruction is what separates 10-tweet threads people read from 20-tweet threads they abandon.
› Create a 30-day content calendar for [PLATFORM] for [BRAND/PERSON]. Content mix: 40% educational, 30% behind-the-scenes/personal, 20% promotional, 10% community/engagement. For each post include: format, topic, hook, and best posting time. Vary formats across the month.	Balanced content strategy that prevents the two failure modes: all promotion (audience leaves) and no promotion (no business results). The format variation prevents content fatigue.
› I need to respond to this [negative comment / complaint / criticism] publicly: [COMMENT]. My response must: acknowledge their experience without admitting fault, show we take it seriously, offer a path to resolution, and read well to the silent majority watching — not just the person complaining. Draft 3 options with different tones.	Public response strategy that addresses two audiences simultaneously — the complainer and the observers. Most brands optimise only for the complainer and inadvertently signal weakness to everyone else watching.

AI IMAGE EDITING & INPAINTING

POWER PROMPT	WHAT IT EFFECTIVELY DOES
› Inpaint this region: [DESCRIPTION OF SELECTED AREA]. The replacement should: match the lighting direction from [side], match the colour temperature of the surrounding area, maintain the same texture style as [reference area], and appear as if it was always part of the original image. Blend distance: [tight/medium/wide].	Photorealistic inpainting with lighting and texture constraints. Lighting direction is the most commonly missed parameter — inpaints that ignore it are immediately visible as AI edits.
› Extend this image [direction: left/right/up/down] by [amount]. The extension should continue: the perspective lines, the lighting gradient, the background texture, and any patterns present. Do not introduce new subjects that weren't implied by the original.	Outpainting with geometric and tonal continuity. The 'no new subjects' constraint prevents the most jarring AI outpainting failure — unexpected objects appearing at the edges.
› Remove [OBJECT/PERSON/TEXT] from this image and fill the background naturally. Reference the texture and pattern from the [surrounding area description] to reconstruct what would realistically be behind the removed element. Maintain consistent depth of field.	Object removal with depth-aware reconstruction. The depth of field constraint prevents the flat, blurry patch that reveals AI removal — the tell that most object removal fails on.
› Apply a style transfer to this image: target style is [REFERENCE STYLE — impressionist, noir, anime, etc.]. Preserve: the subject's identity, spatial composition, and facial features exactly. Transform: colour palette, brushwork/texture, and lighting style. Intensity: [subtle/moderate/full transformation].	Controlled style transfer with identity preservation. The explicit preserve/transform split gives the model precise instructions about what to change and what must stay fixed — the key to style transfer that doesn't destroy the original subject.